



# NaaS: Mirage SDN, Switching and Control

University of Cambridge  
Imperial College London  
University of Nottingham



Richard Mortier, University of Nottingham<sup>W</sup>Cambridge  
<richard.mortier@cl.cam.ac.uk>

# Challenges for NaaS



## 1. Performance & efficiency

- Line rate data processing (10Gbps initially)
- **Efficient use of network resources**
- Offload to hardware when possible

## 2. Programmability & flexibility

- Rapid development of new network services
- Simple deployment, resource placement and allocation

## 3. Security & safety

- Isolation of services within shared hardware

# Current Directions

*A key piece of NaaS is **flexibility**: how can applications better use network resources?*

- OpenFlow Alternatives
  - P<sub>4</sub>, POF
- Raising the abstraction
  - Above the transport layer
- Self-scaling applications
  - Distributed control

# OpenFlow Alternatives

- OpenFlow has limitations:
  - Standard continues to expand (12 to 41 fields)
  - Fixed header fields limits experimentation
  - Extensible match fields provide substrate for extensibility
- Interested in investigating alternatives
  - E.g., **Protocol Oblivious Forwarding** (Huawei), **P4** (Bosshart et al, <http://arxiv.org/pdf/1312.1719.pdf>)

# Emphasising Extensibility

- POF (Huawei)
  - Multi-level table matching
  - (Pointer, Offset), pointer guaranteed to advance
- P<sub>4</sub> (Barefoot, Intel, Stanford, MSR, Google, Princeton)
  - “Parse and populate” model
  - Programmable parser, parallel matching
  - Actions composed of protocol independent switch primitives

*How do these alternatives compare, and are they sufficient for our needs?*

# Raising the Abstraction

- OpenFlow focuses on packets and flows
  - The usual Ethernet and IP headers (incl. IPv6)
  - Plus some support for key control protocols (ARP, DHCP, MPLS, VLAN)
  - Plus statistics messages
- But applications deal in Layer 8 behaviour
  - Mirage *applications* link SDN libraries directly
  - How will applications use them?

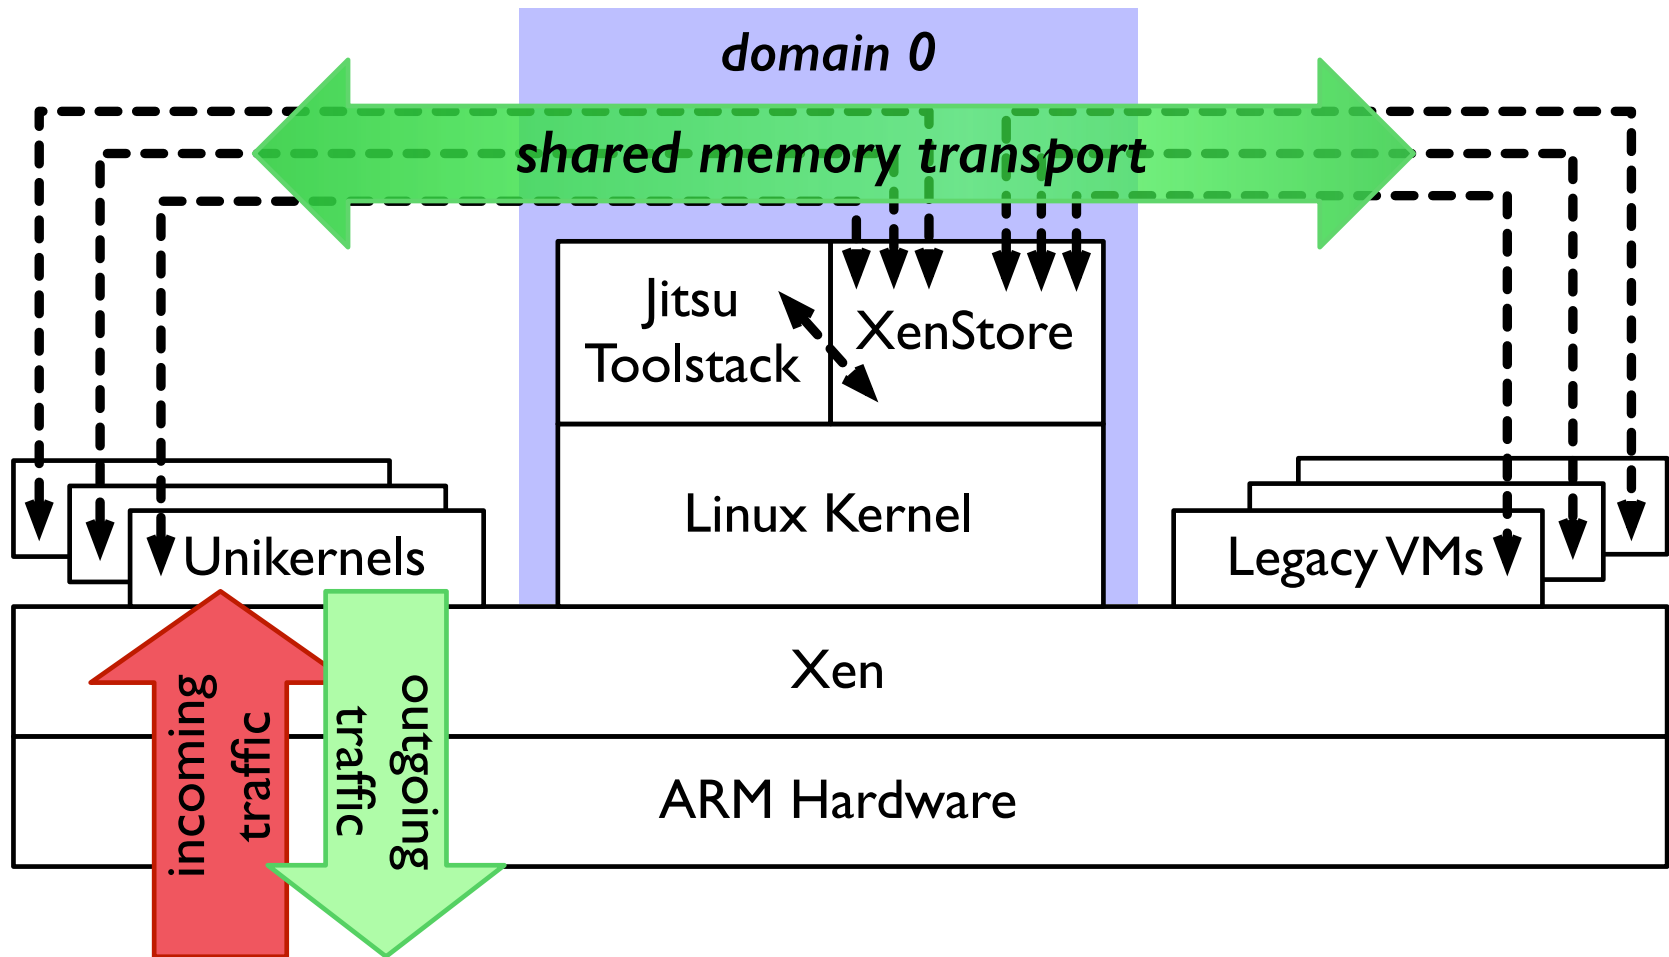
*How can we raise abstractions so that applications can deal in layer 8 concepts (e.g., URLs)?*

# Self-Scaling Applications

- Use of the network by applications is rather naïve
  - The network as a black box
  - Little to no direct control
- SDN allows us to go further
  - Applications install handlers to deal with particular sets of network traffic

*How can we take this further, enabling application logic to modify its own deployment – e.g., scaling up/down in response to load – in light of network conditions?*

# Scaling Out: Jitsu



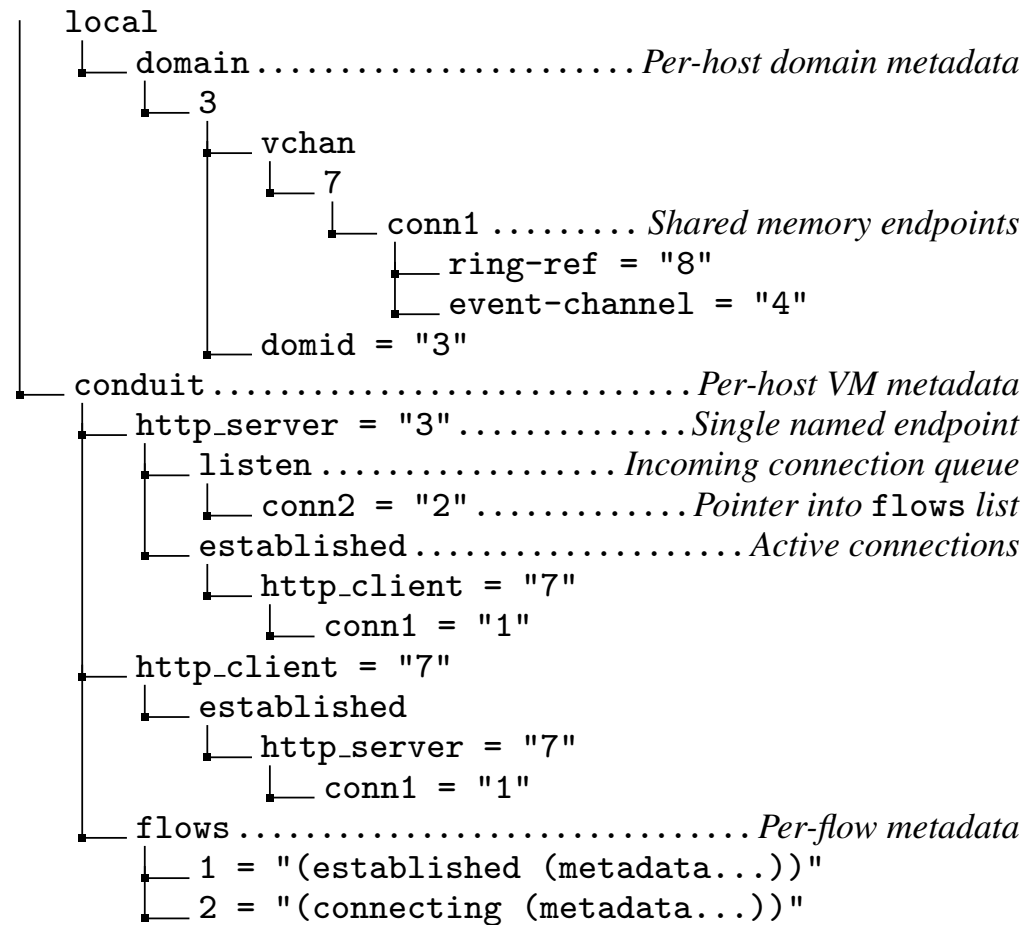


# Conduit: Efficient Inter-VM Comms

Zero-copy shared-memory pages between peers

- Xen grant tables map pages between VMs, synchronised via event channels
  - Rendezvous facility so VMs discover named peers
  - Supports unikernel and legacy VM rendezvous
  - Hooks into higher-level name services like DNS
- Compatible with the *vchan* inter-VM communication protocol

# Conduit: XenStore Layout





# Jitsu Directory Service

- Performs the role of Unix's *inetd*:
  - Jitsu VM launches at boot time to handle name resolution (whether local via a well known jitsud Conduit node in XenStore or remote via DNS)
  - When a request arrives for a live unikernel, Jitsu returns the appropriate endpoint
  - If the unikernel is not live, Jitsu boots it, and acts as proxy until the unikernel is ready

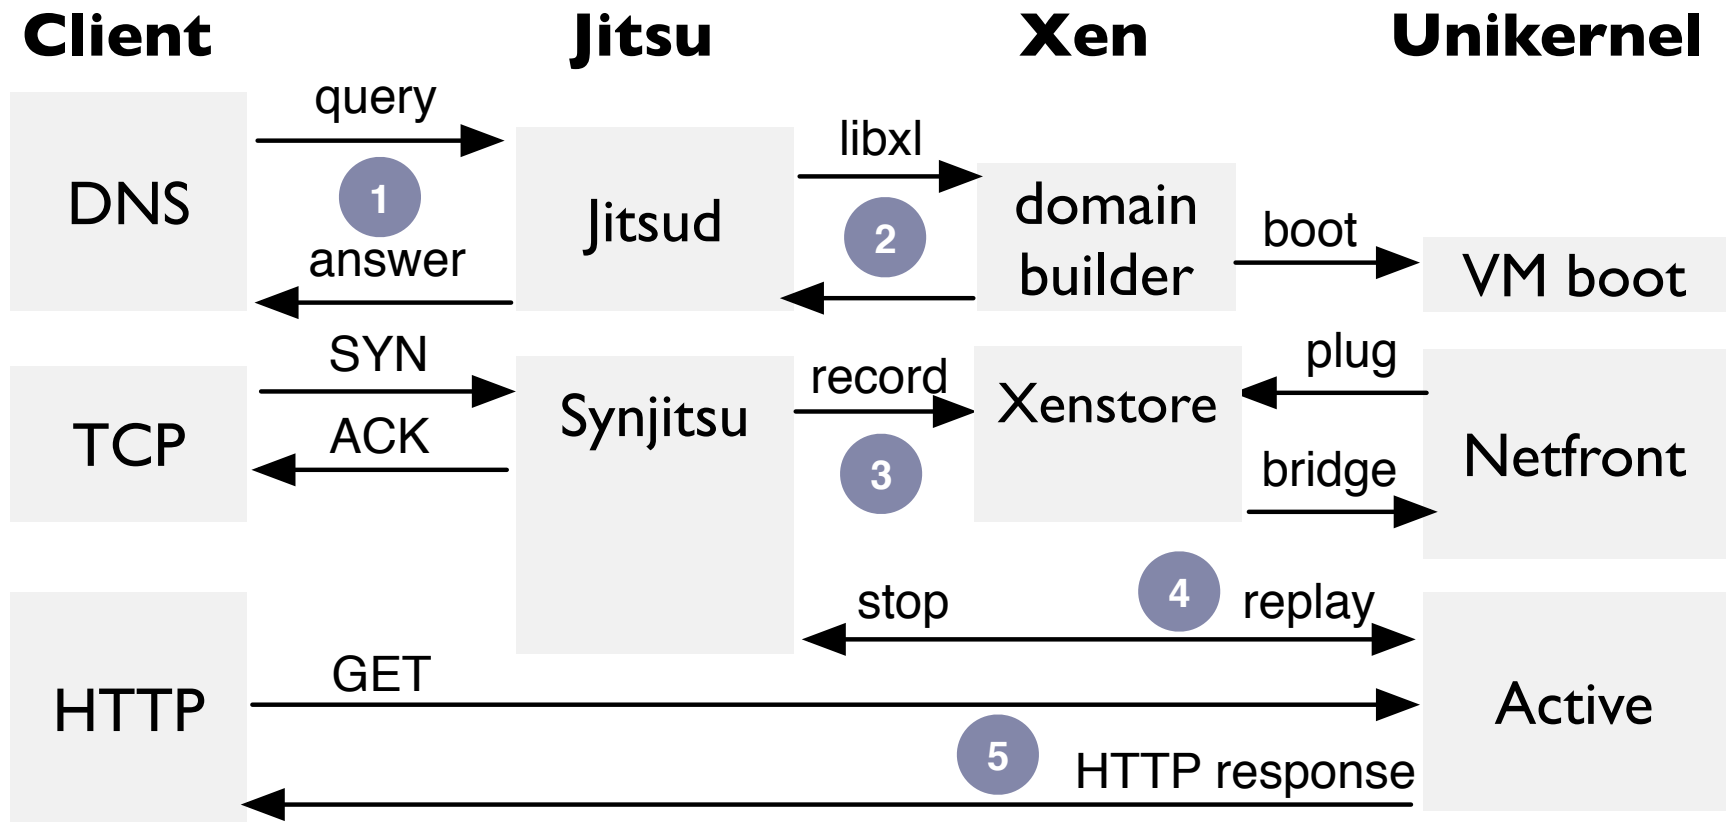
# Result: Low Latency Boot

350ms VM Boot (ARM)



20ms VM Boot (x86)

# Synjitsu: Masking Latency



*By buffering TCP requests into XenStore and then replaying, Synjitsu parallelises connection setup and unikernel boot*

# Summary

So current themes are focused on investigating several ways to improve flexibility

- Replacing OpenFlow
- Raising the level of abstraction
- Automating datacenter deployment